

# TLPTACO v2 – User Guide

**tlptaco** (Teradata Logistics Pipeline for TArgeting and COmmunications) is a modular Python / Jinja-SQL toolkit that turns YAML configuration files into a complete *Eligibility* → *Waterfall* → *Output* campaign pipeline.

This document is the **authoritative, end-to-end “how-to”** for analysts and data engineers who need to configure, run, and extend tlptaco.

## Table of Contents

1. Installation
2. Quick Start
3. Command-Line Interfaces
4. Configuration YAML – full reference
  - Top-level keys
  - Logging
  - Database
  - Eligibility
  - Waterfall
  - Output
  - Failed-records
  - Pre-SQL
5. Engines in depth
  - PreSQLEngine
  - EligibilityEngine
  - WaterfallEngine
  - OutputEngine
6. Working with SQL templates
7. Logging & Diagnostics
8. Appendix A – Example full YAML
9. Appendix B – File/Directory permissions

## Installation

tlptaco is distributed as source; clone the repository and install the Python dependencies into a virtual-environment:

```
git clone <repo>
cd tlptaco
python -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt # includes teradatasql, jinja2, rich, pydantic
```

The project assumes **Python 3.10+** and connectivity to a Teradata database.

## Quick Start

```
# 1 . Write a YAML config (see full reference below)
cp example_yaml.yaml my_campaign.yaml # fill in real values

# 2 . Run the full pipeline (Eligibility → Waterfall → Output)
python -m tlptaco.cli -c my_campaign.yaml -o ./run_artifacts --progress

# 3 . Inspect logs / Excel waterfall / output files in ./run_artifacts

# 4 . Analyse logs afterwards (latest run only)
python -m tlptaco.logs -f run_artifacts/logs/tlptaco.log --summary
```

## Command-Line Interfaces

### 1. Main pipeline runner ( tlptaco.cli )

```
python -m tlptaco.cli -c CONFIG [-o DIR] [--mode full|presizing] [-p] [-v]
```

|                   |   |
|-------------------|---|
| -c / --config     | YAML or JSON config file (required)   |
| -o / --output-dir | Root dir for reports & logs (default: cwd)                                    |
| -m / --mode       | full = Eligibility→Waterfall→Output<br>presizing = Eligibility→Waterfall only |
| -p / --progress   | Rich progress bars  |
| -v / --verbose    | DEBUG console output via rich   |

### 2. Log-inspection utility ( tlptaco.logs )

```
python -m tlptaco.logs -f LOGFILE [-l LEVEL ...] [--summary|--print]
```

```
-l / --levels  DEBUG INFO WARNING ERROR CRITICAL (default = all)
--summary      Show counts per log level (default when no flag specified)
--print        Emit matching lines themselves
```

The parser automatically isolates the **most-recent run** using the header inserted by tlptaco at the start of each run.

## Configuration YAML

All configuration is validated by Pydantic models in `tlptaco.config.schema` – any schema violation aborts execution with a clear message and line/column hint.

Below is a **field-by-field breakdown**. For an "all-in" template jump to [Appendix A](#).

### Top-level keys

| Key              | Type                        | Required | Description  |
|------------------|-----------------------------|----------|--|
| offer_code       | str                         | No       | Short code shown in progress bars & Excel filenames (default <code>Running</code> ). |
| campaign_planner | str                         | No       | Displayed in Excel header.   |
| lead             | str                         | No       | Displayed in Excel header.   |
| logging          | <a href="#">Logging</a>     | Yes      |  |
| database         | <a href="#">Database</a>    | Yes      |  |
| eligibility      | <a href="#">Eligibility</a> | Yes      |  |
| waterfall        | <a href="#">Waterfall</a>   | Yes      |  |
| output           | <a href="#">Output</a>      | Yes      |  |
| pre_sql          | list                        | No       | List of SQL files to run beforehand.   |

### Logging

```
logging:
  level: DEBUG | INFO | WARNING | ERROR | CRITICAL  # console & file
  file:    logs/tlptaco.log          # optional main log
  debug_file: logs/tlptaco.debug.log  # optional DEBUG-only log
  sql_file:  logs/tlptaco.sql.log     # rendered SQL only
  sql_exclude_sections: [waterfall, output] # sections to skip in sql_file
```

Emoji-rich console output is enabled with `-v / --verbose` .

## Database

```
database:
  host: rchtera          # default
  user: my_user
  password: ${TD_PASS}   # env var or plaintext
  logmech: KRB5 | LDAP | TD2 | NONE
```

The DB credentials are consumed by `tlptaco.db.connection.DBConnection` which delegates to the official **teradatasql** driver.

## Eligibility

The engine produces a *smart* table where each **check** becomes a flag column (1 = record passed the check).

```
eligibility:
  eligibility_table: user_wpb.offer123_elig  # schema.table or identifier

# 1. Conditions *****
conditions:
  main:
    BA:          # Baseline checks every record must pass
      - sql: "col1 = 1"          # check 1
        description: "must be active"
      - sql: "date_col >= CURRENT_DATE - 30"

# Per-channel extra checks & segments
channels:
  email:
    BA:
      - sql: "email_opt_in = 1"
      # any key besides BA is treated as *segment*
    loyalty:
      - sql: "loyalty_status = 'GOLD'"

# 2. Tables *****
tables:
  - name: dm.cust_base
    alias: c
    join_type: ''          # omitted = first / base table
    where_conditions: "c.active = 1"

  - name: dm.cust_attrs
    alias: a
```

```
join_type: LEFT
join_conditions: "c.id = a.id"
```

```
# 3. Unique identifiers persisted in final table
unique_identifiers: [c.id, c.household_id]
```

## Key take-aways

- **Name auto-generation** – if you omit the `name:` key under a check it is auto-filled (e.g. `email_loyalty_1` ).
- **Segments** – any sub-mapping under a channel besides `BA` becomes a *segment* that later feeds the waterfall report.

## Waterfall

```
waterfall:
  output_directory: reports/poc/waterfall

  # Records are counted by *groups* - each item can be a single column or
  # a list (composite key). Order matters (rows in Excel).
  count_columns:
    - customer_id
    - [customer_id, account_id]

  history:      # optional historical comparison
    track: true
    recent_window_days: 30      # OR compare_offset_days: 7
    db_path: reports/waterfall_history.sqlite
```

The engine writes one Excel workbook per *run* plus a consolidated workbook  
( `offerCode_YYYY_MM_DD_HH:MM:SS.xlsx` ).

## Output

```
output:
  channels:
    email:
      columns:
        - c.customer_id
        - c.email
        - chk_main1
      column_types:      # optional per-column CASTs
        c.customer_id: "VARCHAR(9)"

    file_location: reports/email
    file_base_name: email_list
```

```

output_options:
  format: csv | excel | parquet | table
  additional_arguments: {sep: "|"} # passed to pandas writer
  custom_function: package.module:transform_df

unique_on: [customer_id] # dedup via QUALIFY ROW_NUMBER()

sms: {...}

# Failed-records optional dump
failed_records:
  enabled: true
  first_reason_only: true
  file_location: reports/failed
  file_base_name: failed_rows
  output_options:
    format: parquet

```

**column\_types** – when provided the output template renders

```
CAST(c.customer_id AS VARCHAR(9)) AS customer_id,
```

otherwise it keeps the column untouched.

## Failed-records dump

When *enabled* the engine writes a flat file containing every record that did **not** pass all checks, annotated with the first failure reason.

## Pre-SQL

```

pre_sql:
- path: sql/prepare_tables.sql
  analytics:
    table: sandbox.tmp_prepare
    unique_counts:
      - customer_id
      - [company_id, site_id]

```

- Each file is split on semicolons ( ; ).
- Optional *analytics* run simple `COUNT(DISTINCT ...)` queries and log the results.

## Engines in depth

## 1. PreSQLEngine

- Reads each `.sql` file, splits on semicolons, executes sequentially via `DBRunner.run()` . • Optional analytics perform quick `COUNT(DISTINCT ...)` and log the counts.

## 2. EligibilityEngine

- Builds a work-table by joining the declared *tables* and computing flag columns for every BA / segment check.
- Post-run it logs total rows and `COUNT(DISTINCT ...)` for each unique identifier.

## 3. WaterfallEngine

- For each `count_columns` group it generates two kinds of SQL (full and segment detail) and pivots the results.
- Excel writer lives in `tlptaco.engines.waterfall_excel` .
- Optional history stored in SQLite enables comparison tabs.

## 4. OutputEngine

- Renders **output.sql.j2** per channel, applies optional `column_types` casting and `CASE` template logic.
- Destination:

- *format = table* → creates a multiset table in Teradata.
- *file formats* → fetches DataFrame then writes CSV / Parquet / Excel via `tlptaco.istream.writer.write_dataframe()` (auto-chmod 770 and `.end` marker file).

# SQL Templates

All templates are Jinja2 files under `tlptaco/sql/templates` . You can list them programmatically:

```
from tlptaco.sql.generator import SQLGenerator
gen = SQLGenerator('tlptaco/sql/templates')
print(gen.list_templates())
```

Fork/extend templates as needed – keep custom ones in a separate directory then instantiate `SQLGenerator(custom_dir)` and monkey-patch the engine in your bootstrap code.

# Logging & Diagnostics

- Main & debug logs are standard text; each run starts with a header:

```
=====
TLPTACO RUN START 2025-08-07 12:00:00
=====
```

- Rendered SQL is saved to a *dedicated* file when `logging.sql_file` is set. Use `sql_exclude_sections` to keep the file lean (e.g. exclude “output”).
- Inspect logs after the fact with the built-in parser:

```
python -m tlptaco.logs -f logs/tlptaco.log --print -l ERROR WARNING
```

## Appendix A – Example full YAML

See `example_yaml.yaml` in the repository – it contains **every** configurable field with inline explanations and choice lists.

## Appendix B – Permissions

When tlptaco writes any file or directory it calls `tlptaco.utils.fs.grant_group_rwx()` which:

1. Tries `os.chown(path, -1, <gid_of_cwd>)` – change *group* ownership to match the working directory.
2. Sets mode `0o770` (rwx for owner & group) so teammates can access the artefacts without manual `chmod/chgrp`.

If the operation fails (platform does not support it or the user lacks privileges) the error is swallowed and execution continues.

Happy targeting! 🎯